

Exercise Sheet 7: Variational Autoencoders

```
import os
import random
import time
import matplotlib.pyplot as plt; plt.rcParamsDefaults()
import numpy as np

# Import Pytorch, Sklearn
from sklearn.manifold import TSNE
import torch
import torch.nn as nn
from torch.nn import functional as F

# Load FashionMNIST dataset
from torchvision import datasets, transforms

# set figure size
plt.figure(figsize=(10, 5))

plt.rcParams['figure.dpi'] = 150
plt.rcParams['axes.titlesize'] = 8
plt.rcParams['axes.labelsize'] = 8
plt.rcParams['xtick.labelsize'] = 6
plt.rcParams['ytick.labelsize'] = 6

<Figure size 1000x500 with 0 Axes>

# set random seed
random.seed(42)
torch.manual_seed(42)

# Set the device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")

#####
##### HOUSEKEEPING #####

# set hyperparameters
EPOCHS = 7
BATCH_SIZE = 128
LEARNING_RATE = 1e-3
WEIGHT_DECAY = 1e-5

# set directories
RESULTS_DIR = 'results/'
DATA_DIR = 'data/'
MODEL_DIR = 'models/'
```

```
#####
#####

# if don't exist, create directories
if not os.path.exists(RESULTS_DIR):
    os.makedirs(RESULTS_DIR)

if not os.path.exists(DATA_DIR):
    os.makedirs(DATA_DIR)

if not os.path.exists(MODEL_DIR):
    os.makedirs(MODEL_DIR)

Using device: cpu

import torch
import torch.nn as nn

class VAE(nn.Module):

    def __init__(self, num_channels=1, num_classes=10, latent_dim=2,
embed_dim=16):
        super(VAE, self).__init__()

        self.latent_dim = latent_dim
        self.embedding = nn.Embedding(num_embeddings=num_classes,
embedding_dim=embed_dim)

        # Encoder
        self.encoder = nn.ModuleList([
            nn.Conv2d(in_channels=num_channels, out_channels=8,
kernel_size=3, stride=2, padding=1),
            nn.Conv2d(in_channels=8, out_channels=16, kernel_size=3,
stride=2, padding=1),
            nn.Conv2d(in_channels=16, out_channels=32, kernel_size=3,
stride=2, padding=1),
        ])

        # Decoder
        self.decoder = nn.ModuleList([
            nn.Conv2d(in_channels=32, out_channels=16, kernel_size=3,
padding=1),
            nn.Conv2d(in_channels=16, out_channels=8, kernel_size=3),
            nn.Conv2d(in_channels=8, out_channels=num_channels,
kernel_size=3, padding=1),
        ])

        # Fully connected layers for learning representations
        self.fc_latent = nn.Linear(in_features=latent_dim + embed_dim,
out_features=512)
```

```

        self.fc_mean = nn.Linear(in_features=512 + embed_dim,
out_features=latent_dim)
        self.fc_var = nn.Linear(in_features=512 + embed_dim,
out_features=latent_dim)
        self.leaky_relu = nn.LeakyReLU()
        self.sigmoid = nn.Sigmoid()

    def forward(self, x, y):
        """
        Args:
            x (tensor): Image(s) of shape [B, C, H, W].
            y (tensor): Class label(s) of shape [B,].

        Returns:
            x_recon (tensor): Reconstructed image(s) of shape [B, C,
H, W].
            mean (tensor): Mean of shape [B, latent_dim].
            log_var (tensor): Log variance of shape [B, latent_dim].
        """
        mean, log_var = self.encode(x, y)
        # Reparameterization Trick
        eps = torch.randn(log_var.shape, device=log_var.device)
        z = mean + torch.exp(log_var * 0.5) * eps
        x_recon = self.decode(z, y)
        return x_recon, mean, log_var

    def encode(self, x, y):
        """
        Args:
            x (tensor): Image(s) of shape [B, C, H, W].
            y (tensor): Class label(s) of shape [B,].

        Returns:
            mean (tensor): Mean of shape [B, latent_dim].
            log_var (tensor): Log variance of shape [B, latent_dim].
        """
        for layer in self.encoder:
            x = layer(x)
            x = self.leaky_relu(x)
        x = torch.reshape(x, (x.shape[0], -1))
        class_embed = self.embedding(y)
        # Concat class information
        mean = self.fc_mean(torch.cat((x, class_embed), dim=1))
        log_var = self.fc_var(torch.cat((x, class_embed), dim=1))
        return mean, log_var

    def decode(self, z, y):
        """
        Args:
            z (tensor): Latent variable(s) of shape [B, latent_dim].

```

```

        y (tensor): Class label(s) of shape [B,].

    Returns:
        x (tensor): Reconstructed image(s) of shape [B, C, H, W].
    """
    class_embed = self.embedding(y)
    # Concat class information
    x = self.fc_latent(torch.cat((z, class_embed), dim=1))
    x = torch.reshape(x, (-1, 32, 4, 4))
    for layer in self.decoder:
        x = nn.functional.interpolate(x, scale_factor=2,
mode='bilinear', align_corners=True)
        x = self.leaky_relu(x)
        x = layer(x)
    x = self.sigmoid(x)
    return x

def sample(self, y, device):
    """
    Args:
        y (int): Class label.
        device (torch.device): Which device to use (cuda or cpu).

    Returns:
        (tensor): Image of shape [1, C, H, W].
    """
    z = torch.randn((1, self.latent_dim), device=device)
    return self.decode(z, torch.tensor([y], device=device))

def sample_latent(self, x, y):
    """
    Args:
        x (tensor): Image(s) of shape [B, C, H, W].
        y (tensor): Class label(s) of shape [B,].

    Returns:
        z (tensor): Latent variable(s) of shape [B, latent_dim].
    """
    mean, log_var = self.encode(x, y)
    # Reparameterization Trick
    eps = torch.randn(log_var.shape, device=log_var.device)
    z = mean + torch.exp(log_var * 0.5) * eps
    return z

```

Task 1: Training a Variational Autoencoder on MNIST

```
# Task 1.1 Download the MNIST dataset
```

```
# Define a transform to normalize the data
```

```
transform = transforms.Compose([transforms.ToTensor()])
```

```
# Download and load the training dataset
```

```
trainset = datasets.MNIST(DATA_DIR, download=True, train=True,  
transform=transform)  
trainloader = torch.utils.data.DataLoader(trainset,  
batch_size=BATCH_SIZE, shuffle=True)
```

```
# Download and load the test dataset
```

```
testset = datasets.MNIST(DATA_DIR, download=True, train=False,  
transform=transform)  
testloader = torch.utils.data.DataLoader(testset,  
batch_size=BATCH_SIZE, shuffle=False)
```

```
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-  
ubyte.gz
```

```
Failed to download (trying next):  
HTTP Error 403: Forbidden
```

```
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-  
images-idx3-ubyte.gz
```

```
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-  
images-idx3-ubyte.gz to data/MNIST/raw/train-images-idx3-ubyte.gz
```

```
100%|██████████| 9912422/9912422 [00:05<00:00, 1770704.65it/s]
```

```
Extracting data/MNIST/raw/train-images-idx3-ubyte.gz to data/MNIST/raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-  
ubyte.gz
```

```
Failed to download (trying next):  
HTTP Error 403: Forbidden
```

```
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-  
labels-idx1-ubyte.gz
```

```
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-  
labels-idx1-ubyte.gz to data/MNIST/raw/train-labels-idx1-ubyte.gz
```

```
100%|██████████| 28881/28881 [00:00<00:00, 237235.82it/s]
```

```
Extracting data/MNIST/raw/train-labels-idx1-ubyte.gz to data/MNIST/raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz  
Failed to download (trying next):  
HTTP Error 403: Forbidden
```

```
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-  
idx3-ubyte.gz
```

```
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-  
idx3-ubyte.gz to data/MNIST/raw/t10k-images-idx3-ubyte.gz
```

```
100%|██████████| 1648877/1648877 [00:00<00:00, 1957087.00it/s]
```

```
Extracting data/MNIST/raw/t10k-images-idx3-ubyte.gz to data/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Failed to download (trying next):
HTTP Error 403: Forbidden
```

```
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz to data/MNIST/raw/t10k-labels-idx1-ubyte.gz
```

```
100%|██████████| 4542/4542 [00:00<00:00, 2137386.82it/s]
```

```
Extracting data/MNIST/raw/t10k-labels-idx1-ubyte.gz to data/MNIST/raw
```

```
num_images = 5

# Plot the images from trainloader
plt.figure(figsize=(num_images, 2))
plt.title('MNIST Dataset Images')
# place it on GPU if available
dataiter = iter(trainloader)
for i in range(num_images):
    # Get the next batch of images
    images, _ = next(dataiter)

    # Plot the first image in the batch
    plt.subplot(1, num_images, i + 1)
    plt.imshow(images[0].numpy().squeeze(), cmap='gray')
    plt.axis('off')

plt.show()
plt.savefig(RESULTS_DIR + 'mnist_dataset_images.png')
```



<Figure size 960x720 with 0 Axes>

```

# Task 1.1 Write pytorch code for loss function (BCE + KL Divergence)
#  $L(x(i), \theta, \phi) = -\mathbb{E}_z \sim q\phi(z|x(i)) [\log p\theta(x(i)|z)] + DKL(q\phi(z|x(i)) || p\theta(z))$ 

def loss_function(x, x_recon, mean, log_var, KL_weight=1e-3):
    """
    Args:
        x (tensor): Image(s) of shape [B, C, H, W].
        x_recon (tensor): Reconstructed image(s) of shape [B, C, H, W].
        mean (tensor): Mean of shape [B, latent_dim].
        log_var (tensor): Log variance of shape [B, latent_dim].

    Returns:
        loss (tensor): Loss value.
    """
    # Ensure the input and output are of the same shape
    assert x.shape == x_recon.shape, "x and x_recon must have the same shape"

    # BCE Loss -- reconstruction loss
    BCE_loss = nn.functional.binary_cross_entropy(x_recon, x, reduction='sum')
    # KL Divergence -- regularization loss
    KL_loss = -0.5 * torch.sum(1 + log_var - mean.pow(2) - log_var.exp())
    # Total loss -- balance between the two (lower KL weight for better reconstruction)
    return BCE_loss + (KL_weight * KL_loss)

def reconstruction_loss_function(x, x_recon):
    """
    Args:
        x (tensor): Image(s) of shape [B, C, H, W].
        x_recon (tensor): Reconstructed image(s) of shape [B, C, H, W].
        mean (tensor): Mean of shape [B, latent_dim].
        log_var (tensor): Log variance of shape [B, latent_dim].

    Returns:
        loss (tensor): Loss value.
    """
    # Ensure the input and output are of the same shape
    assert x.shape == x_recon.shape, "x and x_recon must have the same shape"

    # BCE Loss -- reconstruction loss
    BCE_loss = nn.functional.binary_cross_entropy(x_recon, x, reduction='sum')
    return BCE_loss

```

```
# Task 1.1 Write pytorch code for loss function (BCE + KL Divergence)
#  $L(x(i), \theta, \phi) = -E_z \sim q\phi(z|x(i)) [\log p\theta(x(i)|z)] + DKL(q\phi(z|x(i)) || p\theta(z))$ 
```

```
def regularization_loss_function(mean, log_var):
```

```
    """
```

```
    Args:
```

```
        x (tensor): Image(s) of shape [B, C, H, W].
```

```
        x_recon (tensor): Reconstructed image(s) of shape [B, C, H, W].
```

```
        mean (tensor): Mean of shape [B, latent_dim].
```

```
        log_var (tensor): Log variance of shape [B, latent_dim].
```

```
    Returns:
```

```
        loss (tensor): Loss value.
```

```
    """
```

```
    KL_loss = -0.5 * torch.sum(1 + log_var - mean.pow(2) -
log_var.exp())
```

```
    return KL_loss
```

```
# Task 1.2 Implement the VAE model
```

```
vae = VAE(latent_dim=2).to(device)
```

```
# Implement the ADAM optimizer
```

```
optimizer = torch.optim.Adam(vae.parameters(), lr=LEARNING_RATE,
weight_decay=WEIGHT_DECAY)
```

```
def train(model, optimizer, epochs, trainloader, device,
loss_type='total_loss'):
```

```
    """
```

```
    Train VAE model and plot sample image per class after each epoch.
```

```
    Args:
```

```
        model (nn.Module): VAE model.
```

```
        optimizer (torch.optim): Optimizer.
```

```
        epochs (int): Number of epochs.
```

```
        trainloader (DataLoader): Training data loader.
```

```
        device (torch.device): Which device to use (cuda or cpu).
```

```
    Returns:
```

```
        None
```

```
    """
```

```
    training_losses = []
```

```
    model.train()
```

```
    for epoch in range(epochs):
```

```
        train_loss = 0
```

```
        for i, (images, labels) in enumerate(trainloader):
```

```
            x = images.to(device)
```

```
            y = labels.to(device)
```



```

    # Forward pass
    optimizer.zero_grad()
    x_recon, mean, log_var = model(x, y)

    # Select the loss function
    if loss_type == 'total_loss':
        loss = loss_function(x, x_recon, mean, log_var)
    elif loss_type == 'reconstruction_loss':
        loss = reconstruction_loss_function(x, x_recon)
    elif loss_type == 'regularization_loss':
        loss = regularization_loss_function(mean, log_var)
    else:
        raise ValueError("Invalid loss type. Choose from
['total_loss', 'reconstruction_loss', 'regularization_loss']")

    # Backward pass
    loss.backward()
    optimizer.step()

    train_loss += loss.item()

    # Save the training loss
    training_losses.append(train_loss / len(trainloader.dataset))

    # Print the stats
    print(f'Epoch {epoch + 1}/{epochs}, Training Loss: {train_loss
/ len(trainloader.dataset):.6f}')

    # Plot sample image per class
    sample_images = []
    for i in range(10):
        sample_images.append(model.sample(i,
device).cpu().detach().numpy())

    plt.figure(figsize=(10, 1))
    plt.suptitle(f'Epoch {epoch + 1}/{epochs} - Sample Images')
    for i in range(10):
        plt.subplot(1, 10, i + 1)
        plt.imshow(sample_images[i].squeeze(), cmap='gray')
        plt.axis('off')
    plt.show()
    plt.savefig(RESULTS_DIR + f'sample_images_epoch_{epoch +
1}.png')


    return training_losses

# train the VAE model
training_losses = train(vae, optimizer, EPOCHS, trainloader, device)

```

Epoch 1/7, Training Loss: 194.003004


Epoch 1/7 - Sample Images



Epoch 2/7, Training Loss: 138.144506

<Figure size 960x720 with 0 Axes>


Epoch 2/7 - Sample Images



Epoch 3/7, Training Loss: 134.567883

<Figure size 960x720 with 0 Axes>


Epoch 3/7 - Sample Images



Epoch 4/7, Training Loss: 132.909833

<Figure size 960x720 with 0 Axes>


Epoch 4/7 - Sample Images



Epoch 5/7, Training Loss: 131.950459

<Figure size 960x720 with 0 Axes>


Epoch 5/7 - Sample Images



Epoch 6/7, Training Loss: 131.232240

<Figure size 960x720 with 0 Axes>


Epoch 6/7 - Sample Images



Epoch 7/7, Training Loss: 130.682452

<Figure size 960x720 with 0 Axes>

Epoch 7/7 - Sample Images




<Figure size 960x720 with 0 Axes>

Task 1.3 Plot the training curve (loss) of the VAE model with epochs on x-axis and loss on y-axis

```
def plot_loss_curve(training_losses, figsize=(5, 3),
file_name='training_loss_curve'):
    """ Plot the training loss curve."""
    plt.figure(figsize=figsize)
    plt.plot(training_losses)
    plt.title('Training Loss Curve')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.tight_layout()
    plt.show()
    plt.savefig(RESULTS_DIR + file_name + '.png')
```

Plot the training curve

```
plot_loss_curve(training_losses, figsize=(5, 3),
file_name='vae_training_loss_curve')
```



<Figure size 960x720 with 0 Axes>

Task 2: Visualize the latent space


```
# Task 2.1 train VAE again using only reconstruction loss
vae_rec = VAE(latent_dim=2).to(device)

# Implement the ADAM optimizer
optimizer = torch.optim.Adam(vae_rec.parameters(), lr=LEARNING_RATE,
weight_decay=WEIGHT_DECAY)

# train the VAE model with only reconstruction loss
training_losses = train(vae_rec, optimizer, EPOCHS, trainloader,
device, loss_type='reconstruction_loss')
```

Epoch 1/7, Training Loss: 186.668833


Epoch 1/7 - Sample Images



Epoch 2/7, Training Loss: 138.399085

<Figure size 960x720 with 0 Axes>


Epoch 2/7 - Sample Images



Epoch 3/7, Training Loss: 134.174906

<Figure size 960x720 with 0 Axes>

Epoch 3/7 - Sample Images



Epoch 4/7, Training Loss: 132.291108

<Figure size 960x720 with 0 Axes>

Epoch 4/7 - Sample Images



Epoch 5/7, Training Loss: 131.302706

<Figure size 960x720 with 0 Axes>

Epoch 5/7 - Sample Images



Epoch 6/7, Training Loss: 130.609699

<Figure size 960x720 with 0 Axes>

Epoch 6/7 - Sample Images



Epoch 7/7, Training Loss: 130.181550


<Figure size 960x720 with 0 Axes>

Epoch 7/7 - Sample Images



<Figure size 960x720 with 0 Axes>

```
# Plot the training curve
plot_loss_curve(training_losses, figsize=(5, 3),
file_name='vae_rec_training_loss_curve')
```



<Figure size 960x720 with 0 Axes>

```
# Task 2.2 train VAE again using only regularization loss
```

```
vae_reg = VAE(latent_dim=2).to(device)
```

```
# Implement the ADAM optimizer
```

```
optimizer = torch.optim.Adam(vae_reg.parameters(), lr=LEARNING_RATE,
weight_decay=WEIGHT_DECAY)
```

```
# Train the model
```

```
training_losses = train(vae_reg, optimizer, EPOCHS, trainloader,
device, loss_type='regularization_loss')
```

Epoch 1/7, Training Loss: 0.000381


Epoch 1/7 - Sample Images



Epoch 2/7, Training Loss: 0.000001

<Figure size 960x720 with 0 Axes>


Epoch 2/7 - Sample Images



Epoch 3/7, Training Loss: 0.000000

<Figure size 960x720 with 0 Axes>


Epoch 3/7 - Sample Images



Epoch 4/7, Training Loss: 0.000000

<Figure size 960x720 with 0 Axes>


Epoch 4/7 - Sample Images



Epoch 5/7, Training Loss: 0.000000

<Figure size 960x720 with 0 Axes>

Epoch 5/7 - Sample Images



Epoch 6/7, Training Loss: 0.000000

<Figure size 960x720 with 0 Axes>


Epoch 6/7 - Sample Images



Epoch 7/7, Training Loss: 0.000000


<Figure size 960x720 with 0 Axes>

Epoch 7/7 - Sample Images



<Figure size 960x720 with 0 Axes>

```
# Plot the training curve
plot_loss_curve(training_losses, figsize=(5, 3),
file_name='vae_reg_training_loss_curve')
```



<Figure size 960x720 with 0 Axes>


```
# Task 2.3 train VAE again using both reconstruction and
regularization loss
vae_t = VAE(latent_dim=2).to(device)

# Implement the ADAM optimizer
optimizer = torch.optim.Adam(vae_t.parameters(), lr=LEARNING_RATE,
weight_decay=WEIGHT_DECAY)

# Train the model
training_losses = train(vae_t, optimizer, EPOCHS, trainloader, device)

Epoch 1/7, Training Loss: 195.791238
```

Epoch 1/7 - Sample Images



Epoch 2/7, Training Loss: 137.516372

<Figure size 960x720 with 0 Axes>


Epoch 2/7 - Sample Images



Epoch 3/7, Training Loss: 133.515400

<Figure size 960x720 with 0 Axes>


Epoch 3/7 - Sample Images



Epoch 4/7, Training Loss: 131.846455

<Figure size 960x720 with 0 Axes>


Epoch 4/7 - Sample Images



Epoch 5/7, Training Loss: 130.893433

<Figure size 960x720 with 0 Axes>


Epoch 5/7 - Sample Images



Epoch 6/7, Training Loss: 130.347532


<Figure size 960x720 with 0 Axes>

Epoch 6/7 - Sample Images




Epoch 7/7, Training Loss: 129.867220

<Figure size 960x720 with 0 Axes>



<Figure size 960x720 with 0 Axes>

```
plot_loss_curve(training_losses, figsize=(5, 3),
file_name='vae_total_training_loss_curve')
```



<Figure size 960x720 with 0 Axes>

```
def plot_latent_space(vae, dataloader, n=1000, figsize=5,
device=torch.device('cuda' if torch.cuda.is_available() else 'cpu'),
file_name='vae_latent_space'):
    """ Posterior sampling based on a given MNIST test image. Extract
    the latent space and plot it using t-SNE.

    Args:
        vae (nn.Module): VAE model.
        dataloader (DataLoader): Dataloader to use for sampling from
        the true data distribution.
        n (int): Number of samples to plot.
        figsize (int): Size of the figure.
        method (str): Method to use for projecting the latent space to
        2D. Should be 'TSNE'.
        device (torch.device): Device to use for tensor operations.

    Returns:
```

```

        None
    """
    vae.eval()


    # Sample a large number of points from the true data distribution
    and encode to obtain latent space samples
    latents = []
    labels = []
    with torch.no_grad():
        for i, (data, label) in enumerate(dataloader):
            if len(latents) >= n:
                break
            data = data.to(device)
            label = label.to(device)
            latent = vae.sample_latent(data, label)
            latents.append(latent)
            labels.append(label)
        latents = torch.cat(latents, dim=0)
        labels = torch.cat(labels, dim=0)

    # Project the latent space to 2D using t-SNE
    tsne = TSNE(n_components=2, random_state=42)
    latents_2d = tsne.fit_transform(latents.cpu().numpy())


    # Plot the latent space
    plt.figure(figsize=(figsize, figsize))
    plt.xlabel('z1')
    plt.ylabel('z2')
    plt.title('Latent Space')
    for i in range(10):
        plt.scatter(latents_2d[labels == i, 0], latents_2d[labels ==
i, 1], label=str(i), s=1)
    plt.legend()
    plt.grid()
    plt.tight_layout()
    plt.show()

    # Visualize the latent space
    plot_latent_space(vae, testloader, n=1000, figsize=5, device=device,
file_name='vae_latent_space')


```



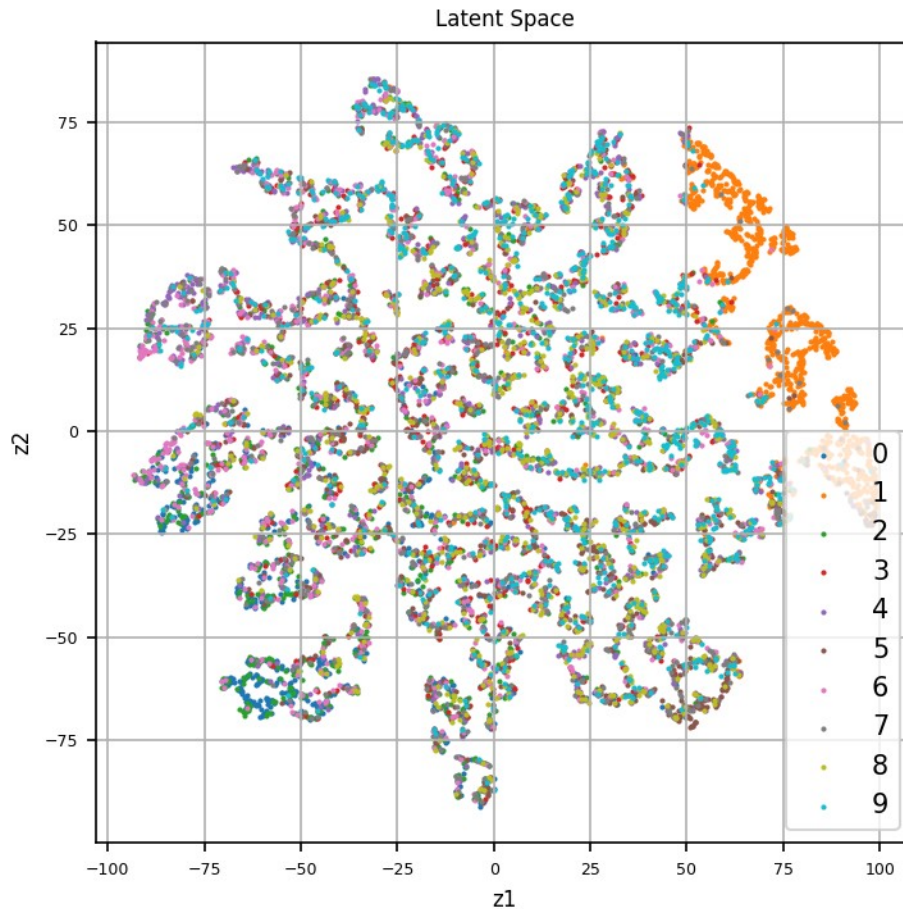
```
# Visualize the latent space for all three models -- reconstruction  
loss only  
plot_latent_space(vae_rec, testloader, n=1000, figsize=5,  
device=device, file_name='vae_rec_latent_space')
```



```
# Visualize the latent space for all three models -- regularization  
loss only  
plot_latent_space(vae_reg, testloader, n=1000, figsize=5,  
device=device, file_name='vae_reg_latent_space')
```



```
# Visualize the latent space -- total loss
plot_latent_space(vae_t, testloader, n=1000, figsize=5, device=device,
file_name='vae_t_latent_space')
```



```
# Plot a grid of samples produced by fixing one latent variable and
varying the other

def plot_latent_traversal(vae, device, n=10):
    """ Plot a grid of samples produced by fixing one latent variable
    and varying the other.

    Args:
        vae (nn.Module): VAE model.
        device (torch.device): Device to use for tensor operations.
        n (int): Number of samples to plot.

    Returns:
        None
    """
    vae.eval()

    # Sample a latent variable based on a standard normal distribution
    z = torch.randn(1, vae.latent_dim), device=device)
    z = z.repeat(n, 1)
```

```


# Plot 10x10 grid of samples
for i in range(0, n):
    # Select fixed label
    y = torch.tensor([i] * n, device=device)

    # Decode the latent variable
    with torch.no_grad():
        samples = vae.decode(z, y)


    # Plot the samples
    plt.figure(figsize=(n, 1))
    for i in range(n):
        plt.subplot(1, n, i + 1)
        plt.imshow(samples[i].squeeze().cpu().numpy(),
cmap='gray')
        plt.axis('off')
        # add class label
        plt.title(f'Class: {y[i].item()}')
    plt.show()
    plt.savefig(RESULTS_DIR + f'vae_latent_traversal_{i}.png')

# Plot the latent traversal
plot_latent_traversal(vae, device, n=10)


```




<Figure size 960x720 with 0 Axes>



<Figure size 960x720 with 0 Axes>




<Figure size 960x720 with 0 Axes>



<Figure size 960x720 with 0 Axes>




<Figure size 960x720 with 0 Axes>




<Figure size 960x720 with 0 Axes>




<Figure size 960x720 with 0 Axes>



<Figure size 960x720 with 0 Axes>



<Figure size 960x720 with 0 Axes>



<Figure size 960x720 with 0 Axes>

Task 3: Anomaly Detection using a Variational Autoencoder

```
# Task 3.1: Download the Fashion MNIST dataset
transform = transforms.Compose([transforms.ToTensor()])


fashion_trainset = datasets.FashionMNIST(DATA_DIR, download=True,
train=True, transform=transform)
fashion_trainloader = torch.utils.data.DataLoader(fashion_trainset,
batch_size=BATCH_SIZE, shuffle=True)

fashion_testset = datasets.FashionMNIST(DATA_DIR, download=True,
train=False, transform=transform)
fashion_testloader = torch.utils.data.DataLoader(fashion_testset,
batch_size=BATCH_SIZE, shuffle=False)

# Plot 5 sample images from trainloader
plt.figure(figsize=(num_images, 2))
plt.title('Fashion MNIST Dataset Images')

dataiter = iter(fashion_trainloader)
for i in range(num_images):
    images, _ = next(dataiter)

    plt.subplot(1, num_images, i + 1)
    plt.imshow(images[0].numpy().squeeze(), cmap='gray')
    plt.axis('off')
```



```
def compare_reconstruction_loss(vae, fashion_testloader,
mnist_testloader, device, file_name='reconstruction_loss_comparison'):
```

```
""" Anomaly detection identifies data points that deviate
significantly from the norm
Compare the reconstruction loss between MNIST and Fashion
MNIST datasets.
```

```
Args:
```

```
vae (nn.Module): VAE model.
trainloader (DataLoader): Training data loader.
testloader (DataLoader): Test data loader.
device (torch.device): Which device to use (cuda or cpu).
```

```
Returns:
```

```
None
```

```
"""
```

```
vae.eval()
```

```
imgs_mnist = []
```

```
imgs_fashion = []
```

```
rec_loss_mnist = []
```

```
rec_loss_fashion = []
```

```
# Compute the reconstruction loss for MNIST dataset
```

```
for i, (images, labels) in enumerate(mnist_testloader):
```

```
    images = images.to(device)
```

```
    labels = labels.to(device)
```

```
    x_recon, mean, log_var = vae(images, labels)
```

```
    loss = reconstruction_loss(images, x_recon)
```

```
    rec_loss_mnist.append(loss.item())
```

```
    imgs_mnist.append(images.cpu().detach().numpy())
```

```
# Compute the reconstruction loss for Fashion MNIST dataset
```

```
for i, (images, labels) in enumerate(fashion_testloader):
```

```
    images = images.to(device)
```

```
    labels = labels.to(device)
```

```
    x_recon, mean, log_var = vae(images, labels)
```

```
    loss = reconstruction_loss(images, x_recon)
```

```
    rec_loss_fashion.append(loss.item())
```

```
    imgs_fashion.append(images.cpu().detach().numpy())
```

```
# Plot 10 sample MNIST images and their corresponding
reconstruction loss values
```

```
fig, axs = plt.subplots(2, figsize=(10, 5))
```

```
fig.suptitle('MNIST Dataset Images and Reconstruction Loss')
```

```
# Bar chart of reconstruction loss
```

```
axs[0].bar(range(1, 11), rec_loss_mnist[:10])
```

```
axs[0].set_xlabel('Image Index')
```

```
axs[0].set_ylabel('Reconstruction Loss')
```

```
axs[0].set_xticks(range(1, 11))
```

```
# set range for y-axis
```

```
axs[0].set_ylim([0, 120000])
```

```

    # Plot a line on top of the bar chart in red connecting each chart
    element
    axs[0].plot(range(1, 11), rec_loss_mnist[:10], color='red',
linewidth=2)

    # MNIST images and their corresponding reconstruction loss values
    for i in range(10):
        ax = plt.subplot(1, 10, i + 1) # create a subplot for each
image
        ax.imshow(imgs_mnist[i][0].squeeze(), cmap='gray')
        ax.text(0.5, -0.1, f'Rec. Loss: {rec_loss_mnist[i]:.2f}',
ha='center', fontsize=4, color='w', transform=ax.transAxes) # display
the reconstruction loss value below the image
        ax.axis('off')

    plt.savefig(f'results/{file_name}_mnist.png')
    plt.show()

    # Plot 10 sample Fashion MNIST images and their corresponding
reconstruction loss values
    fig, axs = plt.subplots(2, figsize=(10, 5))
    fig.suptitle('Fashion MNIST Dataset Images and Reconstruction
Loss')

    # Bar chart of reconstruction loss
    axs[0].bar(range(1, 11), rec_loss_fashion[:10])
    axs[0].set_xlabel('Image Index')
    axs[0].set_ylabel('Reconstruction Loss')
    axs[0].set_xticks(range(1, 11))
    # set range for y-axis
    axs[0].set_ylim([0, 120000])

    # Plot a line on top of the bar chart in red connecting each chart
    element
    axs[0].plot(range(1, 11), rec_loss_fashion[:10], color='red',
linewidth=2)


    # Fashion MNIST images and their corresponding reconstruction loss
    values
    for i in range(10):
        ax = plt.subplot(1, 10, i + 1) # create a subplot for each
image
        ax.imshow(imgs_fashion[i][0].squeeze(), cmap='gray')
        ax.text(0.5, -0.1, f'Rec. Loss: {rec_loss_fashion[i]:.2f}',
ha='center', fontsize=4, color='w', transform=ax.transAxes)
        ax.axis('off')

    plt.savefig(f'results/{file_name}_fashion_mnist.png')
    plt.show()

```

```
# Compare the reconstruction loss between MNIST and Fashion MNIST
# datasets
# Use samples from MNIST handwritten digit dataset and the MNIST.
# fashion dataset to showcase the difference in reconstruction loss.
# Plot 10 samples of each dataset and their corresponding reconstruction
# loss.
compare_reconstruction_loss(vae, fashion_testloader, testloader,
device, file_name='reconstruction_loss_comparison')
```

MNIST Dataset Images and Reconstruction Loss



Fashion MNIST Dataset Images and Reconstruction Loss

